

Fahad_st125981_ML_assignment_3

October 5, 2025

0.0.1 A3: Predicting Car Price

Name: Muhammad Fahad Waqar Student No: 125981

0.0.2 Github Repository Link

https://github.com/mfahadwaqar/st125981_ML_A3

```
[37]: # Importing Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
import warnings
warnings.filterwarnings("ignore")
```

```
[38]: # Loading Data
df = pd.read_csv("Cars.csv") # Reading data from the csv file
print(df.head())
```

		name	year	selling_price	km_driven	fuel	\
0		Maruti Swift Dzire VDI	2014	450000	145500	Diesel	
1		Skoda Rapid 1.5 TDI Ambition	2014	370000	120000	Diesel	
2		Honda City 2017-2020 EXi	2006	158000	140000	Petrol	
3		Hyundai i20 Sportz Diesel	2010	225000	127000	Diesel	
4		Maruti Swift VXi BSIII	2007	130000	120000	Petrol	

	seller_type	transmission	owner	mileage	engine	max_power	\
0	Individual	Manual	First Owner	23.4 kmpl	1248 CC	74 bhp	
1	Individual	Manual	Second Owner	21.14 kmpl	1498 CC	103.52 bhp	
2	Individual	Manual	Third Owner	17.7 kmpl	1497 CC	78 bhp	
3	Individual	Manual	First Owner	23.0 kmpl	1396 CC	90 bhp	
4	Individual	Manual	First Owner	16.1 kmpl	1298 CC	88.2 bhp	

	torque	seats
--	--------	-------

```

0          190Nm@ 2000rpm    5.0
1          250Nm@ 1500-2500rpm  5.0
2          12.7@ 2,700(kgm@ rpm)  5.0
3  22.4 kgm at 1750-2750rpm    5.0
4          11.5@ 4,500(kgm@ rpm)  5.0

```

[39]: *# EDA*

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8128 entries, 0 to 8127
Data columns (total 13 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   name            8128 non-null   object
 1   year            8128 non-null   int64
 2   selling_price   8128 non-null   int64
 3   km_driven       8128 non-null   int64
 4   fuel            8128 non-null   object
 5   seller_type     8128 non-null   object
 6   transmission    8128 non-null   object
 7   owner           8128 non-null   object
 8   mileage         7907 non-null   object
 9   engine          7907 non-null   object
10  max_power       7913 non-null   object
11  torque          7906 non-null   object
12  seats           7907 non-null   float64
dtypes: float64(1), int64(3), object(9)
memory usage: 825.6+ KB

```

[40]: *df.shape # Shape of the dataset*

[40]: (8128, 13)

[41]: *print(df.describe()) # Statistical summary of the dataset*

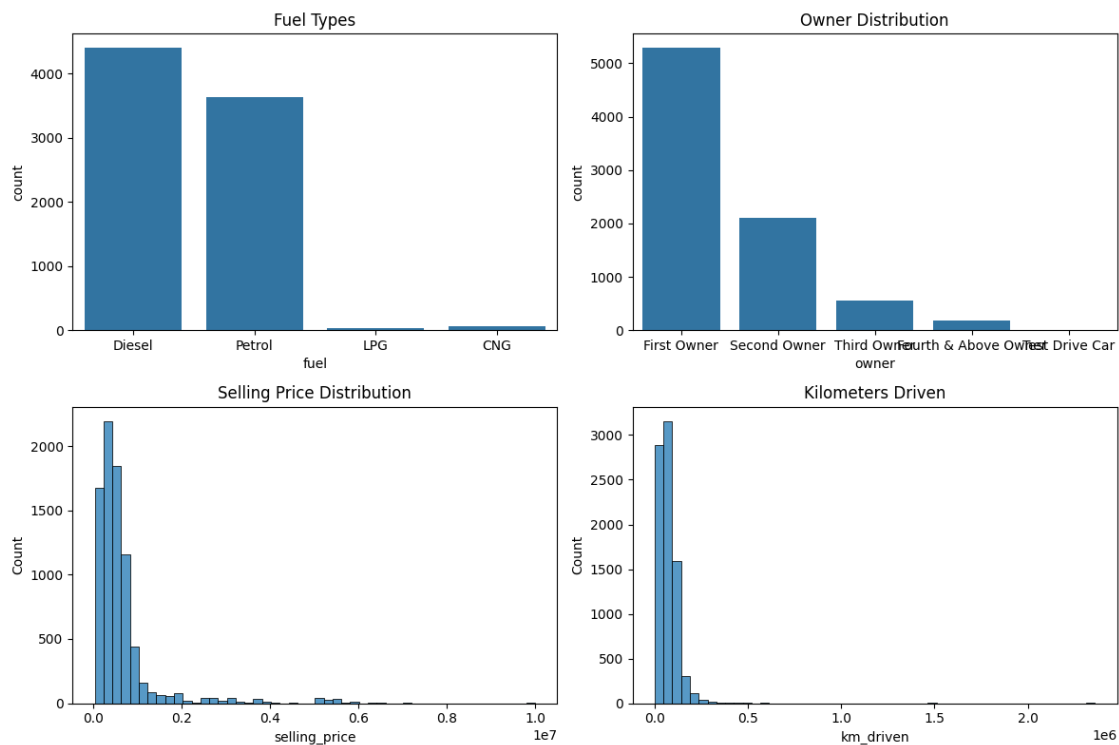
	year	selling_price	km_driven	seats
count	8128.000000	8.128000e+03	8.128000e+03	7907.000000
mean	2013.804011	6.382718e+05	6.981951e+04	5.416719
std	4.044249	8.062534e+05	5.655055e+04	0.959588
min	1983.000000	2.999900e+04	1.000000e+00	2.000000
25%	2011.000000	2.549990e+05	3.500000e+04	5.000000
50%	2015.000000	4.500000e+05	6.000000e+04	5.000000
75%	2017.000000	6.750000e+05	9.800000e+04	5.000000
max	2020.000000	1.000000e+07	2.360457e+06	14.000000

[42]: *# Checking missing values in each column*

```
df.isnull().sum()
```

```
[42]: name          0
      year          0
      selling_price  0
      km_driven     0
      fuel          0
      seller_type   0
      transmission  0
      owner         0
      mileage      221
      engine        221
      max_power     215
      torque        222
      seats         221
      dtype: int64
```

```
[43]: # Univariate Analysis
fig, axes = plt.subplots(2, 2, figsize=(12, 8))
sns.countplot(x="fuel", data=df, ax=axes[0,0]).set_title("Fuel Types")
sns.countplot(x="owner", data=df, ax=axes[0,1]).set_title("Owner Distribution")
sns.histplot(df["selling_price"], bins=50, ax=axes[1,0]).set_title("Selling_
Price Distribution")
sns.histplot(df["km_driven"], bins=50, ax=axes[1,1]).set_title("Kilometers_
Driven")
plt.tight_layout()
```




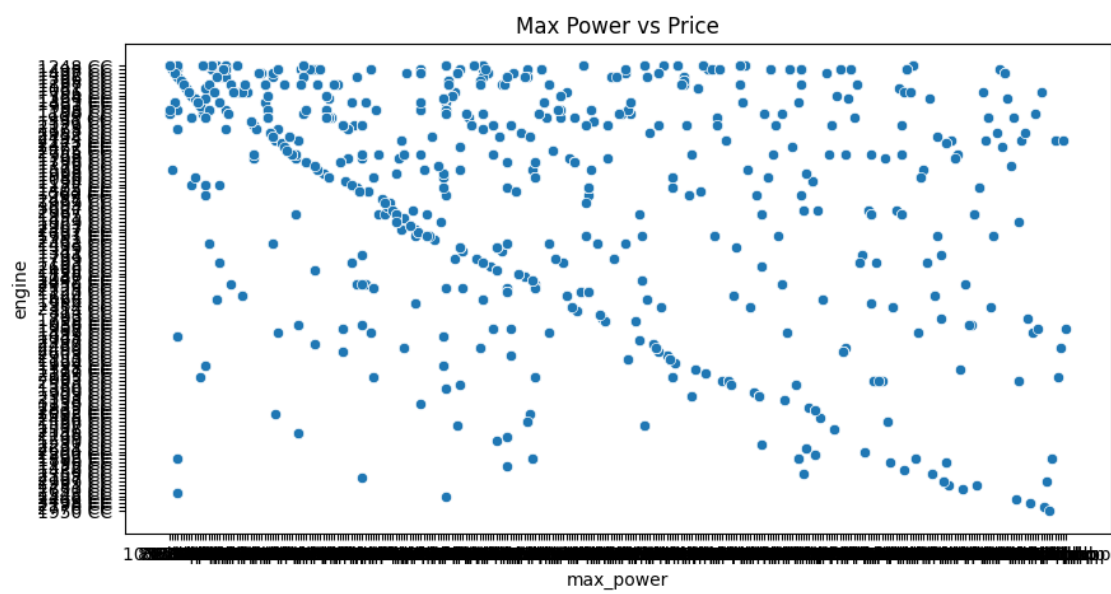
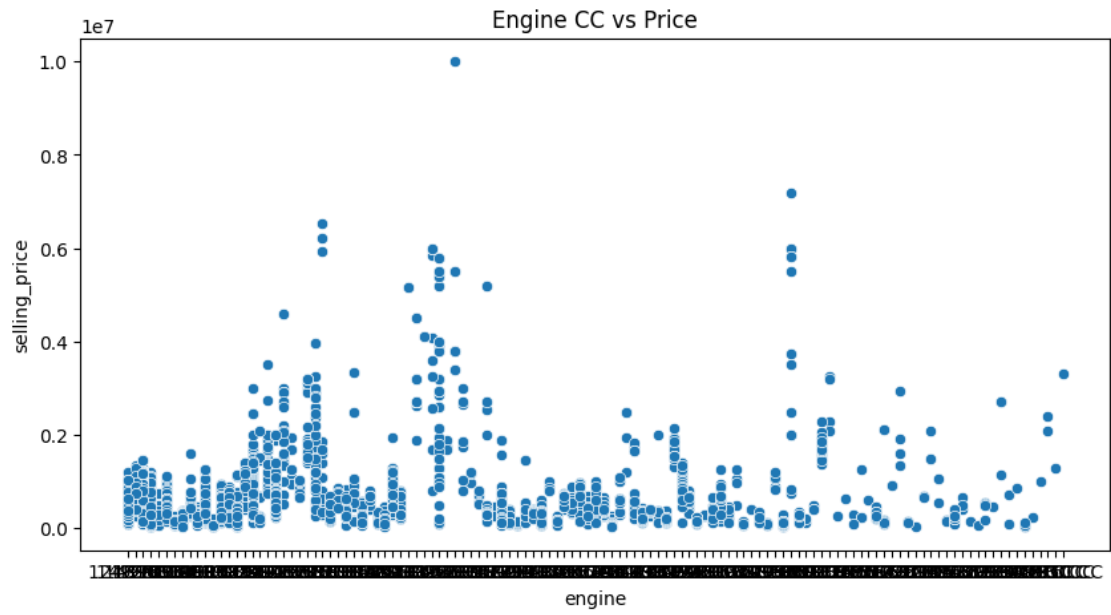
```
[44]: # Multivariate Analysis
plt.figure(figsize=(10, 5))
sns.boxplot(x="owner", y="selling_price", data=df).set_title("Price by Owner_
↪Category")
plt.show()

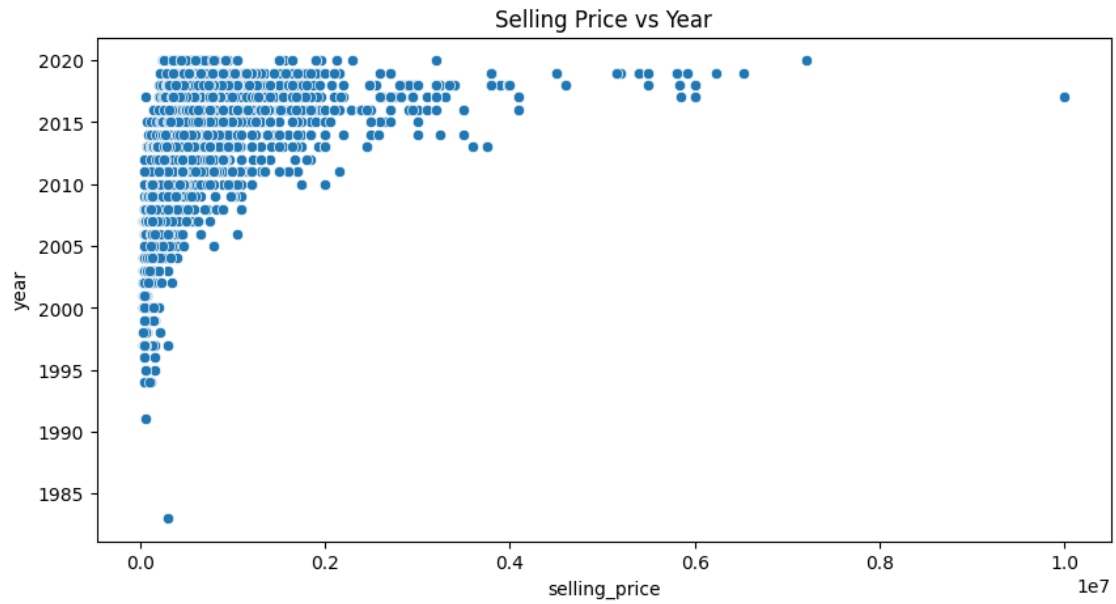
plt.figure(figsize=(10, 5))
sns.scatterplot(x="engine", y="selling_price", data=df).set_title("Engine CC vs_
↪Price")
plt.show()

plt.figure(figsize=(10, 5))
sns.scatterplot(x="max_power", y="engine", data=df).set_title("Max Power vs_
↪Price")
plt.show()

plt.figure(figsize=(10, 5))
sns.scatterplot(x="selling_price", y="year", data=df).set_title("Selling Price_
↪vs Year")
plt.show()
```

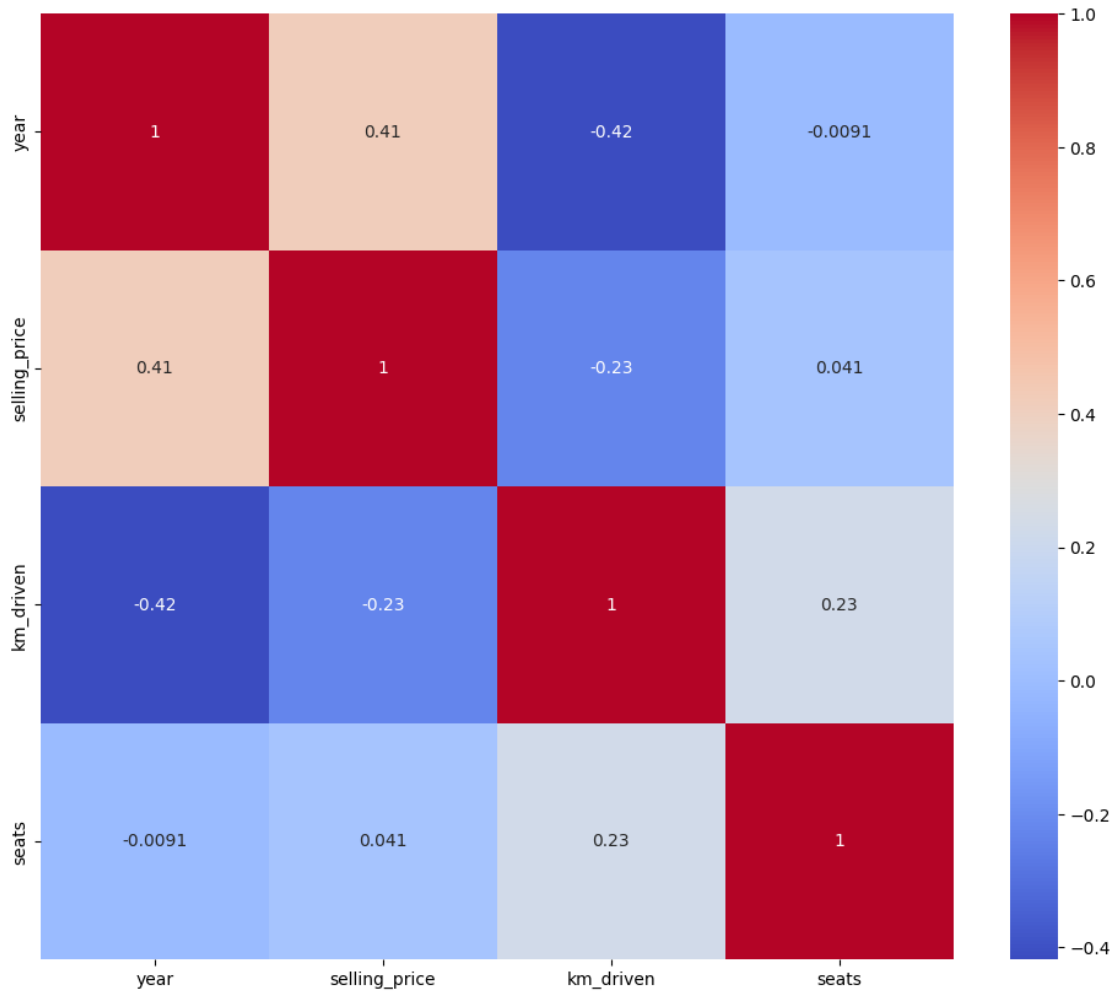






```
[45]: plt.figure(figsize=(12, 10))
sns.heatmap(df.select_dtypes(include='number').corr(), annot=True,
            cmap='coolwarm')
```

```
[45]: <Axes: >
```




```
[46]: import ppscore as pps

matrix_df = pps.matrix(df)[['x', 'y', 'ppscore']].pivot(columns='x', index='y',
↪values='ppscore')

#plot
plt.figure(figsize=(15, 8))
sns.heatmap(matrix_df, vmin=0, vmax=1, cmap="Blues", linewidths=0.5, annot=True)
```

```
[46]: <Axes: xlabel='x', ylabel='y'>
```



```
[47]: # Feature Selection
# Creating a function for data cleaning and feature engineering to be used
# after data splitting
def clean_data(df):
    df = df.copy()

    # Drop torque column
    if "torque" in df.columns:
        df.drop(columns="torque", inplace=True)

    # Remove CNG and LPG
    df = df[~df["fuel"].isin(["CNG", "LPG"])]

    # Remove Test Drive Car from owner
    df = df[df["owner"] != "Test Drive Car"]

    # Convert mileage, engine, max_power to numeric
    df["mileage"] = df["mileage"].str.extract(r'(\d+\.\d*)').astype(float)
    df["engine"] = df["engine"].str.extract(r'(\d+\.\d*)').astype(float)
    df["max_power"] = df["max_power"].str.extract(r'(\d+\.\d*)').astype(float)

    # Extract brand name
    if "name" in df.columns:
        df["brand"] = df["name"].str.split(" ").str[0]
```



```

# Map owner to numeric
owner_mapping = {
    "First Owner": 1,
    "Second Owner": 2,
    "Third Owner": 3,
    "Fourth & Above Owner": 4
}
df["owner"] = df["owner"].map(owner_mapping)

return df

```

```

[48]: # Splitting into train and test sets
train_df, test_df = train_test_split(df, test_size=0.1, random_state=42)

# Applying the cleaning function to both splits
train_df_clean = clean_data(train_df)
test_df_clean = clean_data(test_df)

# Seperating features and targets
X_train = train_df_clean[["year", "km_driven", "mileage", "engine",
    ↪ "max_power", "transmission", "owner", "brand"]]
y_train = np.log(train_df_clean["selling_price"])

X_test = test_df_clean[["year", "km_driven", "mileage", "engine", "max_power",
    ↪ "transmission", "owner", "brand"]]
y_test = np.log(test_df_clean["selling_price"])

print("Train shape:", X_train.shape, "Test shape:", X_test.shape)

```

Train shape: (7229, 8) Test shape: (799, 8)

```

[49]: # Preprocessing
# Rearranging columns
cols = X_train.columns.tolist()
cols.insert(0, cols.pop(cols.index('brand')))
X_train = X_train[cols]

```

```

[50]: cols = X_test.columns.tolist()
cols.insert(0, cols.pop(cols.index('brand')))
X_test = X_test[cols]

```

```

[51]: print(X_train.head())

```

	brand	year	km_driven	mileage	engine	max_power	transmission	\
7287	Mahindra	2019	17000	17.30	1497.0	121.0	Manual	
3027	Maruti	2016	120000	26.59	1248.0	74.0	Manual	
2941	Maruti	2012	100000	19.10	1197.0	85.8	Manual	

4269	BMW	2019	7500	16.78	1995.0	190.0	Automatic
3182	Mahindra	2017	88754	16.00	2179.0	140.0	Manual

	owner
7287	1
3027	3
2941	1
4269	1
3182	1

```
[52]: numeric_features = ["year", "km_driven", "mileage", "engine", "max_power"]
      categorical_features = ["brand", "transmission", "owner"]

      # Split data into numeric and categorical subsets
      X_train_num = X_train[numeric_features].copy()
      X_test_num = X_test[numeric_features].copy()

      X_train_cat = X_train[categorical_features].copy()
      X_test_cat = X_test[categorical_features].copy()

      # Imputation and Scaling for numeric features
      num_medians = X_train_num.median()
      X_train_num = X_train_num.fillna(num_medians)
      X_test_num = X_test_num.fillna(num_medians)

      scaler = StandardScaler()
      X_train_num_scaled = pd.DataFrame(
          scaler.fit_transform(X_train_num),
          columns=numeric_features,
          index=X_train.index
      )
      X_test_num_scaled = pd.DataFrame(
          scaler.transform(X_test_num),
          columns=numeric_features,
          index=X_test.index
      )

      # Imputation and Encoding for categorical features
      cat_modes = X_train_cat.mode().iloc[0]
      X_train_cat = X_train_cat.fillna(cat_modes)
      X_test_cat = X_test_cat.fillna(cat_modes)

      encode = OneHotEncoder(drop="first", sparse_output=False,
                              handle_unknown="ignore")
      X_train_cat_ohe = pd.DataFrame(
          encode.fit_transform(X_train_cat),
          columns=encode.get_feature_names_out(categorical_features),
```

```

        index=X_train.index
    )
X_test_cat_ohe = pd.DataFrame(
    encode.transform(X_test_cat),
    columns=encode.get_feature_names_out(categorical_features),
    index=X_test.index
)

# Putting the processed data back together
X_train = pd.concat([X_train_num_scaled, X_train_cat_ohe], axis=1)
X_test = pd.concat([X_test_num_scaled, X_test_cat_ohe], axis=1)

print("Processed train shape:", X_train.shape)
print("Processed test shape:", X_test.shape)

```

Processed train shape: (7229, 40)

Processed test shape: (799, 40)

0.1 Task 1: Classification & Task 2: Ridge Logistic Regression

```

[53]: # Converting selling_price to 4 discrete classes

# Converting selling_price back from log space to original prices for binning
train_prices_original = np.exp(y_train)
test_prices_original = np.exp(y_test)

# Making 4 bins for classification
bins = pd.qcut(train_prices_original, q=4, labels=[0, 1, 2, 3],
               ↪duplicates='drop')
y_train_class = bins.astype(int)

# Applying the same to test set
bin_edges = pd.qcut(train_prices_original, q=4, duplicates='drop',
                   ↪retbins=True)[1]
y_test_class = pd.cut(test_prices_original, bins=bin_edges, labels=[0, 1, 2,
                           ↪3], include_lowest=True).astype(int)

print(f"Class distribution in training set:")
print(y_train_class.value_counts().sort_index())
print(f"\nClass distribution in test set:")
print(y_test_class.value_counts().sort_index())

```

Class distribution in training set:

selling_price

0 1849

1 1841

2 1743

3 1796

Name: count, dtype: int64

Class distribution in test set:

selling_price

0 201

1 203

2 200

3 195

Name: count, dtype: int64

```
[54]: # Multinomial Logistic Regression Implemetation
class LogisticRegression:
    def __init__(self, learning_rate=0.01, max_iterations=1000,
        ↪regularization=None, lambda_reg=0.01):
        self.learning_rate = learning_rate
        self.max_iterations = max_iterations
        self.regularization = regularization
        self.lambda_reg = lambda_reg
        self.W = None
        self.losses = []

    def softmax(self, theta_t_x):
        theta_t_x = np.array(theta_t_x)
        exp_vals = np.exp(theta_t_x - np.max(theta_t_x, axis=1, keepdims=True))
        return exp_vals / np.sum(exp_vals, axis=1, keepdims=True)

    def h_theta(self, X, W):
        X = np.array(X) # ← Added this line
        return self.softmax(X @ W)

    def gradient(self, X, Y):
        m = X.shape[0]
        h = self.h_theta(X, self.W)

        # Cross-entropy loss
        loss = -np.sum(Y * np.log(h + 1e-8)) / m

        if self.regularization == 'ridge':
            loss += self.lambda_reg * np.sum(self.W ** 2) / (2 * m)

        # Compute gradient
        error = h - Y
        grad = X.T @ error / m

        if self.regularization == 'ridge':
            grad += self.lambda_reg * self.W / m
```

```

        return loss, grad

def fit(self, X, y):
    X = np.array(X)
    y = np.array(y)

    # Get dimensions
    m, n = X.shape
    k = len(np.unique(y)) # number of classes

    # Initialize weights
    self.W = np.random.normal(0, 0.01, (n, k))

    # Convert y to one-hot encoding
    Y_encoded = np.zeros((m, k))
    for i, label in enumerate(y):
        Y_encoded[i, label] = 1

    # Training loop
    self.losses = []
    for i in range(self.max_iterations):
        loss, grad = self.gradient(X, Y_encoded)
        self.losses.append(loss)
        self.W = self.W - self.learning_rate * grad

        if i % 100 == 0:
            print(f"Loss at iteration {i}: {loss}")

def predict(self, X):
    X = np.array(X)
    return np.argmax(self.h_theta(X, self.W), axis=1)

def predict_proba(self, X):
    X = np.array(X)
    return self.h_theta(X, self.W)

def plot_losses(self):
    plt.figure(figsize=(10, 6))
    plt.plot(self.losses)
    plt.title('Training Loss Over Iterations')
    plt.xlabel('Iteration')
    plt.ylabel('Cross-Entropy Loss')
    plt.grid(True)
    plt.show()

```

```

[55]: # Evaluation Metrics
class ClassificationMetrics:

```

```

def __init__(self, y_true, y_pred, classes=None):
    self.y_true = np.array(y_true)
    self.y_pred = np.array(y_pred)
    self.classes = classes if classes is not None else np.unique(np.
→concatenate([y_true, y_pred]))
    self.n_classes = len(self.classes)
    self._compute_confusion_matrix()

def _compute_confusion_matrix(self):
    self.confusion_matrix = np.zeros((self.n_classes, self.n_classes),
→dtype=int)
    for i, true_class in enumerate(self.classes):
        for j, pred_class in enumerate(self.classes):
            self.confusion_matrix[i, j] = np.sum((self.y_true ==
→true_class) & (self.y_pred == pred_class))

def accuracy(self):
    correct_predictions = np.sum(self.y_true == self.y_pred)
    total_predictions = len(self.y_true)
    return correct_predictions / total_predictions

def precision(self, class_idx=None):
    if class_idx is not None:
        tp = self.confusion_matrix[class_idx, class_idx]
        fp = np.sum(self.confusion_matrix[:, class_idx]) - tp
        return tp / (tp + fp) if (tp + fp) > 0 else 0.0
    else:
        precisions = []
        for i in range(self.n_classes):
            precisions.append(self.precision(i))
        return np.array(precisions)

def recall(self, class_idx=None):
    if class_idx is not None:
        tp = self.confusion_matrix[class_idx, class_idx]
        fn = np.sum(self.confusion_matrix[class_idx, :]) - tp
        return tp / (tp + fn) if (tp + fn) > 0 else 0.0
    else:
        recalls = []
        for i in range(self.n_classes):
            recalls.append(self.recall(i))
        return np.array(recalls)

def f1_score(self, class_idx=None):
    if class_idx is not None:
        prec = self.precision(class_idx)
        rec = self.recall(class_idx)

```

```

        return 2 * (prec * rec) / (prec + rec) if (prec + rec) > 0 else 0.0
    else:
        f1_scores = []
        for i in range(self.n_classes):
            f1_scores.append(self.f1_score(i))
        return np.array(f1_scores)

def macro_precision(self):
    precisions = self.precision()
    return np.mean(precisions)

def macro_recall(self):
    recalls = self.recall()
    return np.mean(recalls)

def macro_f1(self):
    f1_scores = self.f1_score()
    return np.mean(f1_scores)

def weighted_precision(self):
    precisions = self.precision()
    support = [np.sum(self.y_true == class_label) for class_label in self.
↪classes]
    total_support = np.sum(support)
    weights = [s / total_support for s in support]
    return np.sum([w * p for w, p in zip(weights, precisions)])

def weighted_recall(self):
    recalls = self.recall()
    support = [np.sum(self.y_true == class_label) for class_label in self.
↪classes]
    total_support = np.sum(support)
    weights = [s / total_support for s in support]
    return np.sum([w * r for w, r in zip(weights, recalls)])

def weighted_f1(self):
    f1_scores = self.f1_score()
    support = [np.sum(self.y_true == class_label) for class_label in self.
↪classes]
    total_support = np.sum(support)
    weights = [s / total_support for s in support]
    return np.sum([w * f for w, f in zip(weights, f1_scores)])

def classification_report(self):
    support = [np.sum(self.y_true == class_label) for class_label in self.
↪classes]

```

```

        print(f"{'Class':<10} {'Precision':<10} {'Recall':<10} {'F1-Score':<10}␣
↪{'Support':<10}")
        print("-" * 50)

        precisions = self.precision()
        recalls = self.recall()
        f1_scores = self.f1_score()

        for i, class_label in enumerate(self.classes):
            print(f"{class_label:<10} {precisions[i]:<10.3f} {recalls[i]:<10.
↪3f} {f1_scores[i]:<10.3f} {support[i]:<10}")

        print("-" * 50)
        print(f"{'Accuracy':<10} {'':<10} {'':<10} {self.accuracy():<10.3f} {np.
↪sum(support):<10}")
        print(f"{'Macro Avg':<10} {self.macro_precision():<10.3f} {self.
↪macro_recall():<10.3f} {self.macro_f1():<10.3f} {np.sum(support):<10}")
        print(f"{'Weighted Avg':<10} {self.weighted_precision():<10.3f} {self.
↪weighted_recall():<10.3f} {self.weighted_f1():<10.3f} {np.sum(support):<10}")

```

```

[56]: # Testing Implementation
      # Training custom model
      print("Training custom Logistic Regression...")
      custom_model = LogisticRegression(learning_rate=0.01, max_iterations=500)
      custom_model.fit(X_train, y_train_class)

      # Make predictions
      y_pred_custom = custom_model.predict(X_test)

      # Evaluate with custom metrics
      print("Custom Implementation Results")
      custom_metrics = ClassificationMetrics(y_test_class, y_pred_custom)
      custom_metrics.classification_report()

      # Plot training loss
      custom_model.plot_losses()

      # Compare with sklearn
      from sklearn.linear_model import LogisticRegression as SklearnLogisticRegression
      from sklearn.metrics import classification_report

      print("Comparing with sklearn")
      sklearn_model = SklearnLogisticRegression(multi_class='multinomial',␣
↪solver='lbfgs', max_iter=500)
      sklearn_model.fit(X_train, y_train_class)
      y_pred_sklearn = sklearn_model.predict(X_test)

```



```
print("Sklearn Classification Report:")
print(classification_report(y_test_class, y_pred_sklearn))
```

Training custom Logistic Regression...

Loss at iteration 0: 1.3840307045241076

Loss at iteration 100: 1.1802860057536837

Loss at iteration 200: 1.084184448588857

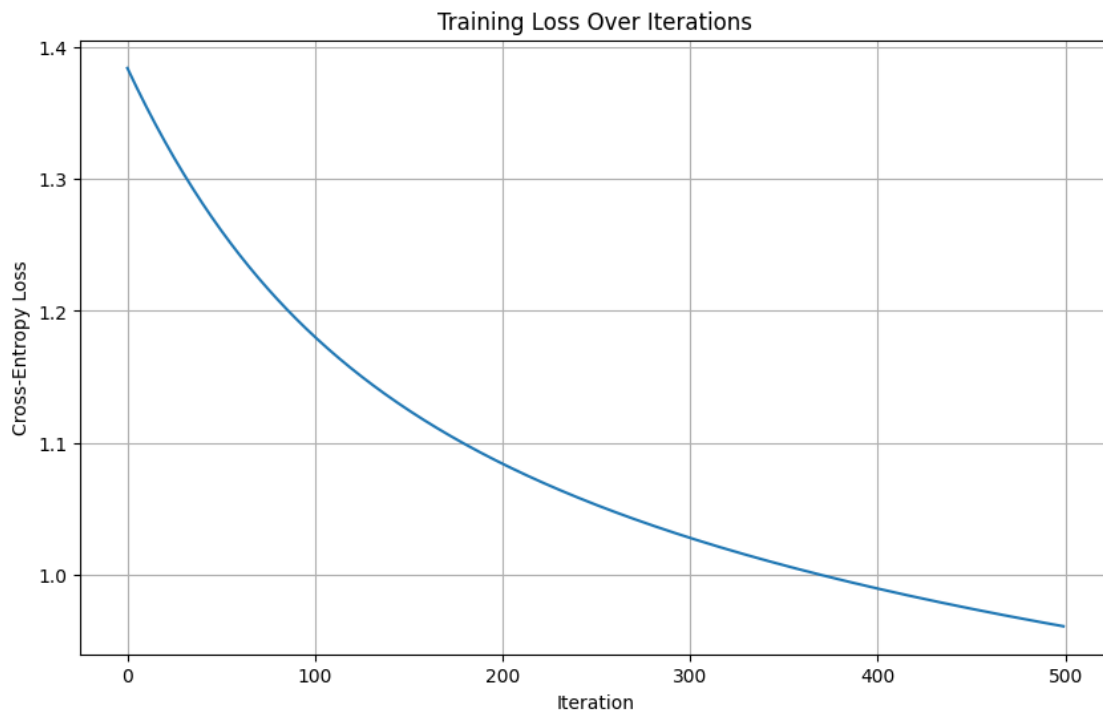
Loss at iteration 300: 1.028002722695251

Loss at iteration 400: 0.9896152504095714

Custom Implementation Results

Class	Precision	Recall	F1-Score	Support
0	0.651	0.935	0.767	201
1	0.557	0.335	0.418	203
2	0.485	0.415	0.447	200
3	0.659	0.733	0.694	195

Accuracy			0.603	799
Macro Avg	0.588	0.605	0.582	799
Weighted Avg	0.588	0.603	0.581	799



Comparing with sklearn

Sklearn Classification Report:

precision	recall	f1-score	support
-----------	--------	----------	---------

0	0.85	0.84	0.85	201
1	0.61	0.60	0.60	203
2	0.55	0.62	0.58	200
3	0.83	0.75	0.79	195
accuracy			0.70	799
macro avg	0.71	0.70	0.70	799
weighted avg	0.71	0.70	0.70	799

```
[57]: # Ridge Logistic Regression Test
print("Ridge Logistic Regression Test")

# Train with Ridge regularization
ridge_model = LogisticRegression(learning_rate=0.01, max_iterations=500,
                                regularization='ridge', lambda_reg=0.1)
ridge_model.fit(X_train, y_train_class)
y_pred_ridge = ridge_model.predict(X_test)

ridge_metrics = ClassificationMetrics(y_test_class, y_pred_ridge)
print("Ridge Regularized Results:")
ridge_metrics.classification_report()
```

```
Ridge Logistic Regression Test
Loss at iteration 0: 1.390120576275421
Loss at iteration 100: 1.1840401718231859
Loss at iteration 200: 1.086821503984475
Loss at iteration 300: 1.0300346672984226
Loss at iteration 400: 0.9912695053227906
Ridge Regularized Results:
```

Class	Precision	Recall	F1-Score	Support
0	0.651	0.935	0.767	201
1	0.558	0.355	0.434	203
2	0.482	0.400	0.437	200
3	0.656	0.723	0.688	195

Accuracy			0.602	799
Macro Avg	0.587	0.603	0.582	799
Weighted Avg	0.586	0.602	0.581	799

0.2 Task 3: Deployment

```
[58]: # MLFlow Experiment Logging
import mlflow
import mlflow.sklearn
from sklearn.model_selection import cross_val_score
import pickle
```

```
import os
```

```
[61]: # Setup MLflow tracking with authentication
mlflow.set_tracking_uri("http://mlflow.ml.brain.cs.ait.ac.th/")

# Set authentication credentials (from A2 instructions: username='admin',
↳password='password')
os.environ['MLFLOW_TRACKING_USERNAME'] = 'admin'
os.environ['MLFLOW_TRACKING_PASSWORD'] = 'password'

mlflow.set_experiment("st125981-a3") # Replace st125981 with your student ID

def log_model_experiment(model, model_name, X_train, X_test, y_train, y_test,
↳params):
    with mlflow.start_run(run_name=model_name):
        # Log parameters
        for key, value in params.items():
            mlflow.log_param(key, value)

        # Log the trained model (so it can be registered later)
        mlflow.sklearn.log_model(model, artifact_path="model")

        # Make predictions
        y_pred = model.predict(X_test)

        # Calculate metrics using your custom implementation
        metrics_calc = ClassificationMetrics(y_test, y_pred)

        # Log all metrics to MLflow
        mlflow.log_metric("accuracy", metrics_calc.accuracy())
        mlflow.log_metric("macro_precision", metrics_calc.macro_precision())
        mlflow.log_metric("macro_recall", metrics_calc.macro_recall())
        mlflow.log_metric("macro_f1", metrics_calc.macro_f1())
        mlflow.log_metric("weighted_precision", metrics_calc.
↳weighted_precision())
        mlflow.log_metric("weighted_recall", metrics_calc.weighted_recall())
        mlflow.log_metric("weighted_f1", metrics_calc.weighted_f1())

        # Log per-class metrics
        precisions = metrics_calc.precision()
        recalls = metrics_calc.recall()
        f1_scores = metrics_calc.f1_score()

        for i, class_label in enumerate(metrics_calc.classes):
            mlflow.log_metric(f"precision_class_{class_label}", precisions[i])
            mlflow.log_metric(f"recall_class_{class_label}", recalls[i])
            mlflow.log_metric(f"f1_class_{class_label}", f1_scores[i])
```

```

        print(f"Logged {model_name} - Run ID: {mlflow.active_run().info.
↳run_id}")
        return mlflow.active_run().info.run_id

```

```

[ ]: # Experiment 1: Basic Logistic Regression
print("Experiment 1: Basic Logistic Regression")
basic_model = LogisticRegression(learning_rate=0.01, max_iterations=500,↳
↳regularization=None)
basic_model.fit(X_train, y_train_class)

basic_params = {
    "model_type": "LogisticRegression",
    "learning_rate": 0.01,
    "max_iterations": 500,
    "regularization": "none",
    "lambda_reg": 0.0
}

basic_run_id = log_model_experiment(
    basic_model, "basic_logistic_regression",
    X_train, X_test, y_train_class, y_test_class,
    basic_params
)

# Experiment 2: Ridge Logistic Regression
print("\nExperiment 2: Ridge Logistic Regression")
ridge_model = LogisticRegression(learning_rate=0.01, max_iterations=500,
                                regularization='ridge', lambda_reg=0.1)
ridge_model.fit(X_train, y_train_class)

ridge_params = {
    "model_type": "LogisticRegression",
    "learning_rate": 0.01,
    "max_iterations": 500,
    "regularization": "ridge",
    "lambda_reg": 0.1
}

ridge_run_id = log_model_experiment(
    ridge_model, "ridge_logistic_regression",
    X_train, X_test, y_train_class, y_test_class,
    ridge_params
)

# Experiment 3: Different Learning Rates
print("\nExperiment 3: Testing Different Learning Rates")

```

```

learning_rates = [0.001, 0.01, 0.1]

for lr in learning_rates:
    print(f"Testing learning_rate={lr}")

    temp_model = LogisticRegression(learning_rate=lr, max_iterations=500)
    temp_model.fit(X_train, y_train_class)

    temp_params = {
        "model_type": "LogisticRegression",
        "learning_rate": lr,
        "max_iterations": 500,
        "regularization": "none",
        "lambda_reg": 0.0
    }

    log_model_experiment(
        temp_model, f"logistic_regression_lr_{lr}",
        X_train, X_test, y_train_class, y_test_class,
        temp_params
    )

# Experiment 4: Different Regularization Strengths
print("\nExperiment 4: Testing Different Regularization Strengths")
lambda_values = [0.01, 0.1, 1.0]

for lam in lambda_values:
    print(f"Testing lambda={lam}")

    temp_model = LogisticRegression(learning_rate=0.01, max_iterations=500,
                                    regularization='ridge', lambda_reg=lam)
    temp_model.fit(X_train, y_train_class)

    temp_params = {
        "model_type": "LogisticRegression",
        "learning_rate": 0.01,
        "max_iterations": 500,
        "regularization": "ridge",
        "lambda_reg": lam
    }

    log_model_experiment(
        temp_model, f"ridge_logistic_regression_lambda_{lam}",
        X_train, X_test, y_train_class, y_test_class,
        temp_params
    )

```

```

# Experiment 5: Different Max Iterations
print("\nExperiment 5: Testing Different Max Iterations")
max_iterations = [100, 500, 1000]

for max_iter in max_iterations:
    print(f"Testing max_iterations={max_iter}")

    temp_model = LogisticRegression(learning_rate=0.01, max_iterations=max_iter)
    temp_model.fit(X_train, y_train_class)

    temp_params = {
        "model_type": "LogisticRegression",
        "learning_rate": 0.01,
        "max_iterations": max_iter,
        "regularization": "none",
        "lambda_reg": 0.0
    }

    log_model_experiment(
        temp_model, f"logistic_regression_maxiter_{max_iter}",
        X_train, X_test, y_train_class, y_test_class,
        temp_params
    )

print(f"View results at: http://mlflow.ml.brain.cs.ait.ac.th/")
print(f"Experiment name: st125981-a3")

print("Finding and registering best model...")

# Get all runs from the experiment
experiment = mlflow.get_experiment_by_name("st125981-a3")
runs = mlflow.search_runs(experiment_ids=[experiment.experiment_id])

# Find the best run based on weighted_f1 score
best_run = runs.loc[runs['metrics.weighted_f1'].idxmax()]
best_run_id = best_run['run_id']
best_weighted_f1 = best_run['metrics.weighted_f1']

print(f"\nBest model found:")
print(f"  Run ID: {best_run_id}")
print(f"  Run Name: {best_run['tags.mlflow.runName']}")
print(f"  Weighted F1: {best_weighted_f1:.4f}")
print(f"  Accuracy: {best_run['metrics.accuracy']:.4f}")

from mlflow import MlflowClient

client = MlflowClient()

```

```

model_name = "st125981_best_logistic_regression"

# Register the best model from its run ID
model_uri = f"runs://{best_run_id}/model"

# Register a new version
model_version = mlflow.register_model(model_uri=model_uri, name=model_name)

# Wait for the registration to complete (optional)
import time
time.sleep(5)

# Assign or update the alias (e.g., "Staging")
client.set_registered_model_alias(model_name, "Staging", model_version.version)

print(f"\nBest model registered successfully!")
print(f"Model Name: {model_name}")
print(f"Model Version: {model_version.version}")
print(f"Alias: Staging")
print(f"MLflow URL: http://mlflow.ml.brain.cs.ait.ac.th/#/models/{model_name}")

```

Experiment 1: Basic Logistic Regression

Loss at iteration 0: 1.3952661230561991
 Loss at iteration 100: 1.1855473260017506
 Loss at iteration 200: 1.0869718985599426
 Loss at iteration 300: 1.0296683778895568
 Loss at iteration 400: 0.990682469274693

2025/10/04 17:35:54 WARNING mlflow.models.model: `artifact_path` is deprecated.
 Please use `name` instead.

2025/10/04 17:35:59 WARNING mlflow.utils.environment: Failed to resolve
 installed pip version. ``pip`` will be added to conda.yaml environment spec
 without a version specifier.

2025/10/04 17:35:59 WARNING mlflow.models.model: Model logged without a
 signature and input example. Please set `input_example` parameter when logging
 the model to auto infer the model signature.

Logged basic_logistic_regression - Run ID: 5b1fecb19ac0473d9d10773e5a9ab000

View run basic_logistic_regression at: <http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407/runs/5b1fecb19ac0473d9d10773e5a9ab000>

View experiment at:
<http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407>

Experiment 2: Ridge Logistic Regression

Loss at iteration 0: 1.384708515079588
 Loss at iteration 100: 1.1821779835479942
 Loss at iteration 200: 1.0862831863873634
 Loss at iteration 300: 1.0299207655820624
 Loss at iteration 400: 0.9912919301327046

2025/10/04 17:36:13 WARNING mlflow.models.model: `artifact_path` is deprecated.
Please use `name` instead.
2025/10/04 17:36:15 WARNING mlflow.utils.environment: Failed to resolve
installed pip version. ``pip`` will be added to conda.yaml environment spec
without a version specifier.
2025/10/04 17:36:15 WARNING mlflow.models.model: Model logged without a
signature and input example. Please set `input_example` parameter when logging
the model to auto infer the model signature.

Logged ridge_logistic_regression - Run ID: 6fd556ffe509429ca77adc4cd351e8c5
View run ridge_logistic_regression at: <http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407/runs/6fd556ffe509429ca77adc4cd351e8c5>
View experiment at:
<http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407>

Experiment 3: Testing Different Learning Rates

Testing learning_rate=0.001
Loss at iteration 0: 1.387765267225824
Loss at iteration 100: 1.3586839106171098
Loss at iteration 200: 1.3320109304683263
Loss at iteration 300: 1.3075514267319492
Loss at iteration 400: 1.2851135089502983

2025/10/04 17:36:29 WARNING mlflow.models.model: `artifact_path` is deprecated.
Please use `name` instead.
2025/10/04 17:36:32 WARNING mlflow.utils.environment: Failed to resolve
installed pip version. ``pip`` will be added to conda.yaml environment spec
without a version specifier.
2025/10/04 17:36:32 WARNING mlflow.models.model: Model logged without a
signature and input example. Please set `input_example` parameter when logging
the model to auto infer the model signature.

Logged logistic_regression_lr_0.001 - Run ID: 288b90b814544ede935c8ea0aa48d995
View run logistic_regression_lr_0.001 at: <http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407/runs/288b90b814544ede935c8ea0aa48d995>
View experiment at:
<http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407>

Testing learning_rate=0.01
Loss at iteration 0: 1.3831223081282826
Loss at iteration 100: 1.181550377466806
Loss at iteration 200: 1.0857468433247324
Loss at iteration 300: 1.0294293709979594
Loss at iteration 400: 0.9908511114753946

2025/10/04 17:36:46 WARNING mlflow.models.model: `artifact_path` is deprecated.
Please use `name` instead.
2025/10/04 17:36:49 WARNING mlflow.utils.environment: Failed to resolve
installed pip version. ``pip`` will be added to conda.yaml environment spec
without a version specifier.
2025/10/04 17:36:49 WARNING mlflow.models.model: Model logged without a

signature and input example. Please set `input_example` parameter when logging the model to auto infer the model signature.

Logged logistic_regression_lr_0.01 - Run ID: aabd2778112c426f9807fcb0ac2ca137

View run logistic_regression_lr_0.01 at: <http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407/runs/aabd2778112c426f9807fcb0ac2ca137>

View experiment at:

<http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407>

Testing learning_rate=0.1

Loss at iteration 0: 1.3907288677231786

Loss at iteration 100: 0.8769298044003523

Loss at iteration 200: 0.8040785011532885

Loss at iteration 300: 0.7675581357690282

Loss at iteration 400: 0.7445452161245624

2025/10/04 17:37:03 WARNING mlflow.models.model: `artifact_path` is deprecated. Please use `name` instead.

2025/10/04 17:37:06 WARNING mlflow.utils.environment: Failed to resolve installed pip version. ``pip`` will be added to conda.yaml environment spec without a version specifier.

2025/10/04 17:37:06 WARNING mlflow.models.model: Model logged without a signature and input example. Please set `input_example` parameter when logging the model to auto infer the model signature.

Logged logistic_regression_lr_0.1 - Run ID: 6fb48e66e798477083264e84b5977dc2

View run logistic_regression_lr_0.1 at: <http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407/runs/6fb48e66e798477083264e84b5977dc2>

View experiment at:

<http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407>

Experiment 4: Testing Different Regularization Strengths

Testing lambda=0.01

Loss at iteration 0: 1.3857680409497024

Loss at iteration 100: 1.1821036186803537

Loss at iteration 200: 1.0858147696811016

Loss at iteration 300: 1.0293042852391872

Loss at iteration 400: 0.9906041614557589

2025/10/04 17:37:20 WARNING mlflow.models.model: `artifact_path` is deprecated. Please use `name` instead.

2025/10/04 17:37:23 WARNING mlflow.utils.environment: Failed to resolve installed pip version. ``pip`` will be added to conda.yaml environment spec without a version specifier.

2025/10/04 17:37:23 WARNING mlflow.models.model: Model logged without a signature and input example. Please set `input_example` parameter when logging the model to auto infer the model signature.

Logged ridge_logistic_regression_lambda_0.01 - Run ID:

e2b82cae170a44778d6463b802122d2b

View run ridge_logistic_regression_lambda_0.01 at: <http://mlflow.ml.brain.cs.a>

it.ac.th/#/experiments/403651241041143407/runs/e2b82cae170a44778d6463b802122d2b

View experiment at:

<http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407>

Testing lambda=0.1

Loss at iteration 0: 1.3909243915430425

Loss at iteration 100: 1.183914052452798

Loss at iteration 200: 1.0862729793372345

Loss at iteration 300: 1.0292647632172647

Loss at iteration 400: 0.9903813120476955

2025/10/04 17:37:37 WARNING mlflow.models.model: `artifact_path` is deprecated.
Please use `name` instead.

2025/10/04 17:37:40 WARNING mlflow.utils.environment: Failed to resolve
installed pip version. ``pip`` will be added to conda.yaml environment spec
without a version specifier.

2025/10/04 17:37:40 WARNING mlflow.models.model: Model logged without a
signature and input example. Please set `input_example` parameter when logging
the model to auto infer the model signature.

Logged ridge_logistic_regression_lambda_0.1 - Run ID:

2204452fb2ee487fb63b5c1e16b7adb1

View run ridge_logistic_regression_lambda_0.1 at: <http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407/runs/2204452fb2ee487fb63b5c1e16b7adb1>

View experiment at:

<http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407>

Testing lambda=1.0

Loss at iteration 0: 1.3832970067721662

Loss at iteration 100: 1.1804134197537357

Loss at iteration 200: 1.0845309351269128

Loss at iteration 300: 1.0283833767481942

Loss at iteration 400: 0.989993338102195

2025/10/04 17:37:54 WARNING mlflow.models.model: `artifact_path` is deprecated.
Please use `name` instead.

2025/10/04 17:37:57 WARNING mlflow.utils.environment: Failed to resolve
installed pip version. ``pip`` will be added to conda.yaml environment spec
without a version specifier.

2025/10/04 17:37:57 WARNING mlflow.models.model: Model logged without a
signature and input example. Please set `input_example` parameter when logging
the model to auto infer the model signature.

Logged ridge_logistic_regression_lambda_1.0 - Run ID:

a06280067883484b8896b7cfd7447845

View run ridge_logistic_regression_lambda_1.0 at: <http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407/runs/a06280067883484b8896b7cfd7447845>

View experiment at:

<http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407>

Experiment 5: Testing Different Max Iterations

Testing max_iterations=100

Loss at iteration 0: 1.3851654098852488

2025/10/04 17:38:11 WARNING mlflow.models.model: `artifact_path` is deprecated.
Please use `name` instead.

2025/10/04 17:38:14 WARNING mlflow.utils.environment: Failed to resolve
installed pip version. ``pip`` will be added to conda.yaml environment spec
without a version specifier.

2025/10/04 17:38:14 WARNING mlflow.models.model: Model logged without a
signature and input example. Please set `input_example` parameter when logging
the model to auto infer the model signature.

Logged logistic_regression_maxiter_100 - Run ID:

277626cc3b9042f69112c8a1b6490bfc

View run logistic_regression_maxiter_100 at: <http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407/runs/277626cc3b9042f69112c8a1b6490bfc>

View experiment at:

<http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407>

Testing max_iterations=500

Loss at iteration 0: 1.3821592825259754

Loss at iteration 100: 1.1800206556650719

Loss at iteration 200: 1.084317934469051

Loss at iteration 300: 1.0281748793083316

Loss at iteration 400: 0.9897495453439152

2025/10/04 17:38:28 WARNING mlflow.models.model: `artifact_path` is deprecated.
Please use `name` instead.

2025/10/04 17:38:31 WARNING mlflow.utils.environment: Failed to resolve
installed pip version. ``pip`` will be added to conda.yaml environment spec
without a version specifier.

2025/10/04 17:38:31 WARNING mlflow.models.model: Model logged without a
signature and input example. Please set `input_example` parameter when logging
the model to auto infer the model signature.

Logged logistic_regression_maxiter_500 - Run ID:

194fa30302324f9295ee23e55f69776f

View run logistic_regression_maxiter_500 at: <http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407/runs/194fa30302324f9295ee23e55f69776f>

View experiment at:

<http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407>

Testing max_iterations=1000

Loss at iteration 0: 1.3872384610129262

Loss at iteration 100: 1.1826853920718026

Loss at iteration 200: 1.0861152579836295

Loss at iteration 300: 1.0295173104317563

Loss at iteration 400: 0.9908016339187147

Loss at iteration 500: 0.9617365013726558

Loss at iteration 600: 0.9386394212600057

Loss at iteration 700: 0.9195930783655573

Loss at iteration 800: 0.9034759555442992

Loss at iteration 900: 0.8895733603072701

2025/10/04 17:38:47 WARNING mlflow.models.model: `artifact_path` is deprecated. Please use `name` instead.

2025/10/04 17:38:50 WARNING mlflow.utils.environment: Failed to resolve installed pip version. ``pip`` will be added to conda.yaml environment spec without a version specifier.

2025/10/04 17:38:50 WARNING mlflow.models.model: Model logged without a signature and input example. Please set `input_example` parameter when logging the model to auto infer the model signature.

Logged logistic_regression_maxiter_1000 - Run ID:
3e85b36d1cfe48c789dbbde28bf42e93

View run logistic_regression_maxiter_1000 at: <http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407/runs/3e85b36d1cfe48c789dbbde28bf42e93>

View experiment at:
<http://mlflow.ml.brain.cs.ait.ac.th/#/experiments/403651241041143407>

View results at: <http://mlflow.ml.brain.cs.ait.ac.th/>

Experiment name: st125981-a3

Finding and registering best model...

Best model found:

Run ID: 6fb48e66e798477083264e84b5977dc2

Run Name: logistic_regression_lr_0.1

Weighted F1: 0.6502

Accuracy: 0.6521

Registered model 'st125981_best_logistic_regression' already exists. Creating a new version of this model...

2025/10/04 17:39:04 WARNING mlflow.tracking._model_registry.fluent: Run with id 6fb48e66e798477083264e84b5977dc2 has no artifacts at artifact path 'model', registering model based on models:/m-ec45281410224c00876c8d6c516b6cfe instead

2025/10/04 17:39:04 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: st125981_best_logistic_regression, version 1

Created version '1' of model 'st125981_best_logistic_regression'.

Best model registered successfully!

Model Name: st125981_best_logistic_regression

Model Version: 1

Alias: Staging

MLflow URL:
http://mlflow.ml.brain.cs.ait.ac.th/#/models/st125981_best_logistic_regression